

MATH 305

Hyperparameter Optimization Using Reinforcement Learning

Alican Akca -Yusuf Erol

Content

1. Methodology

a. Environment

b. Agent

c. Optimization Problem

d. Reward and Penalties

e. Visualization

2. Algorithm

3. Evaluation of the Inferences


4. References

METHODOLOGY

Methodology:

Environment

The goal is to train an agent to balance a pole on a cart by moving the cart left or right.



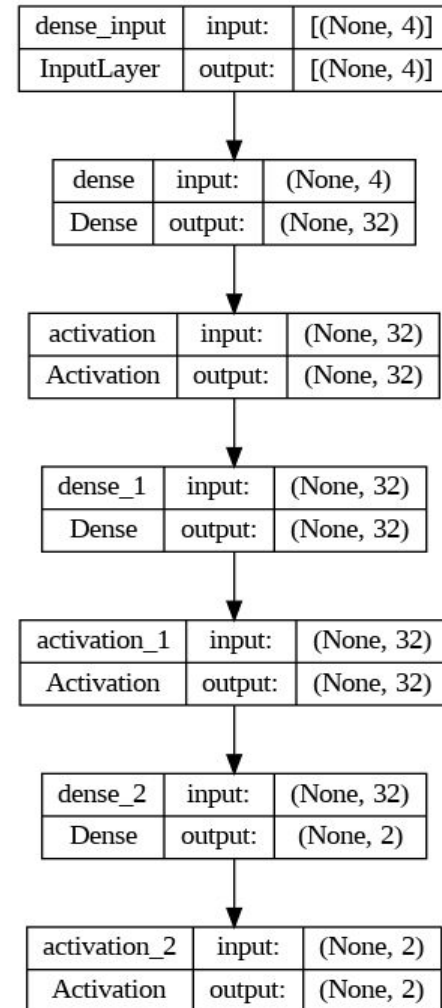
Action Space	Discrete(2)
Observation Shape	(4,)
Observation High	[4.8 inf 0.42 inf]
Observation Low	[-4.8 -inf -0.42 -inf]
Import	<code>gym.make("CartPole-v1")</code>

Methodology:

Agent

The agent makes decisions by choosing actions that are expected to maximize the cumulative reward over time.

The agent is a neural network model that takes the state of the environment as input and outputs the action to be taken.



Methodology:

Optimization Problem

The objective function is to minimize the validation loss (val_loss) which is the Mean Squared Error (MSE).

Given a set of n samples, where for each sample i , the predicted value is \hat{y}_i and the actual value is y_i , the MSE is calculated as:

$$MSE = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

Methodology:

Rewards and Penalties

In the context of reinforcement learning, Q-values that must be maximized represent the expected future reward for taking a certain action in a certain state.

$$Q(s, a) = r + \gamma \max_a Q(s', a')$$

where:

- s is the current state,
- a is the action taken,
- r is the immediate reward received after taking action a in state s ,
- s' is the new state after taking action a ,
- a' is the action taken in state s' ,
- γ is the discount factor which determines the present value of future rewards.

Algorithm

Algorithm

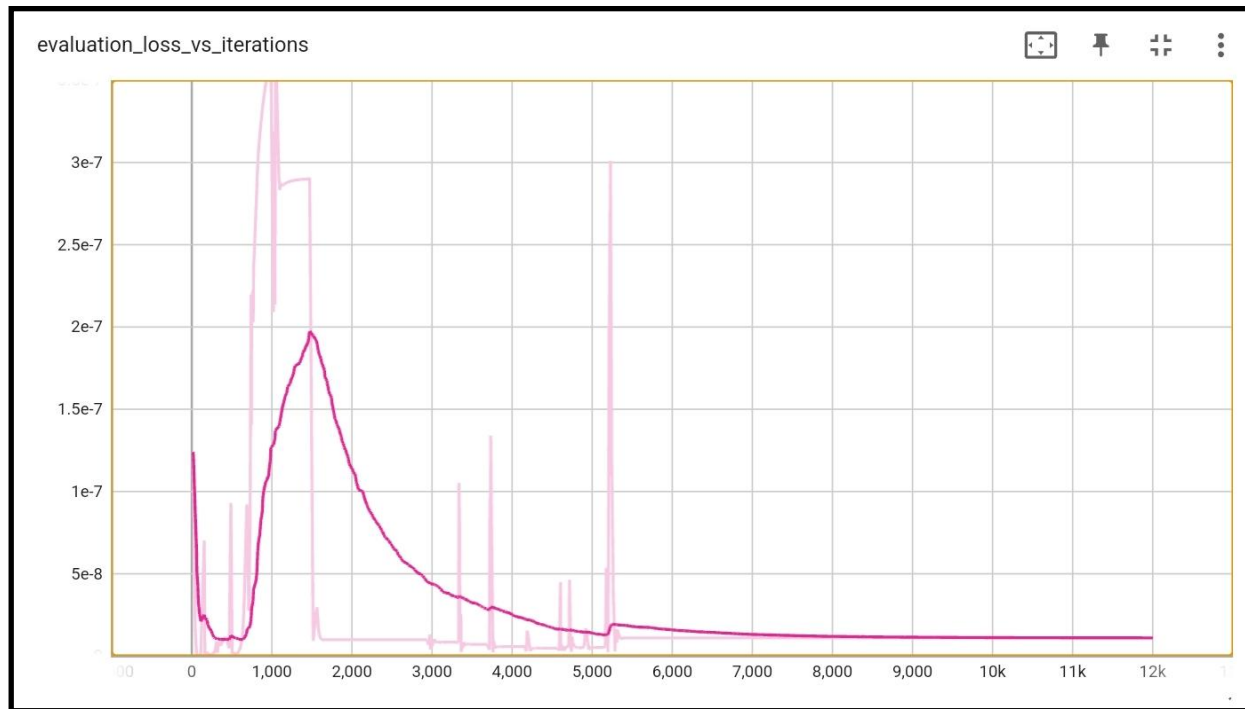
Let $f: \mathbb{R}^n \rightarrow \mathbb{R}$ be the fitness or cost function which must be minimized. Let $\mathbf{x} \in \mathbb{R}^n$ designate a position or candidate solution in the search-space. The basic RandomSearch algorithm can then be described as:

1. Initialize \mathbf{x} with a random position in the search-space.
2. Until a termination criterion is met (e.g. number of iterations performed, or adequate fitness reached), repeat the following:
 1. Sample a new position \mathbf{y} from the hypersphere of a given radius surrounding the current position \mathbf{x} (see e.g. Marsaglia's technique for sampling a hypersphere.)
 2. If $f(\mathbf{y}) < f(\mathbf{x})$ then move to the new position by setting $\mathbf{x} = \mathbf{y}$

Evaluation of the Inferences

Evaluation of the Inferences

- Stability
- Efficiency
- Robustness
- Generalization



References

[1] Neuronlike Adaptive Elements That Can Solve Difficult Learning Control Problems ANDREW G. BARTO, Member, IEEE, Richard S. Sutton, and Charles w. Anderson (0018-9472/83/0900-0834\$01.00 © 1983 IEEE)

Thanks for listening!